

This project aims to predict used car prices in UAE based on specifications and condition

```
In [13]: # Step 1: Load the Data
import pandas as pd

# Load the dataset
df = pd.read_csv('UAE_Used_cars.csv', encoding='UTF-8-SIG')

# Display the first few rows of the dataframe
print(df.head(), df.describe())
```

	Car Brand	Car Model	Production Year	Mileage	Price	\
0	Nissan	Altima	2005	445,740 km	3,500	
1	Toyota	Camry	1999	200,000 km	5,500	
2	Ford	Focus	2006	366,135 km	5,500	
3	Toyota	Echo	2005	200,000 km	6,000	
4	Chevrolet	Epica	2009	250,000 km	6000	
	Description			Specs		
0	Dubai			GCC Specs		
1	Perfect Condition Toyota Camry			GCC Specs		
2	FORD FOCUS			GCC Specs		
3	GCC - TOYOTA ECHO 2005 - Manual, Urgent Sale			GCC Specs		
4	Chevrolet Epica			American Specs		
	Timestamp	Location		Production Year		
0	04-03-24 14:49	Dubai				
1	04-03-24 14:49	Dubai				
2	04-03-24 14:49	Dubai				
3	04-03-24 14:49	Dubai				
4	45354.94097	Abu Dhabi				
count	8006.000000					
mean	2017.939046					
std	5.227208					
min	1929.000000					
25%	2015.000000					
50%	2019.000000					
75%	2022.000000					
max	2024.000000					



```
In [18]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load the data
df = pd.read_csv('UAE_Used_cars.csv', encoding='UTF-8-SIG')

# Convert Price to numeric, removing commas and converting to float
df['Price'] = df['Price'].str.replace(',', '').astype(float)

# Sort by Price in descending order
df_sorted = df.sort_values('Price', ascending=False)

# Display the top 20 cars
print(df_sorted[['Car Brand', 'Car Model', 'Production Year', 'Price', 'Mileage']].head(
```

	Car Brand	Car Model	Production Year	Price	Mileage
8005	Ford	GT	2022	4750000.0	30 km
8004	Ford	GT	2022	3900000.0	30 km
8003	Ford	GT	2021	2880000.0	75 km
8002	Ford	GT	2020	2649000.0	4,000 km
8001	Ford	Mustang	1967	1600000.0	1,749 km
8000	Nissan	GT-R	1999	949999.0	121,454 km
7999	Nissan	GT-R	1999	949999.0	137,488 km
7998	Ford	Mustang	1968	799999.0	43,200 km
7997	Nissan	Patrol	2013	750000.0	230,000 km
7996	Nissan	GT-R	2022	599000.0	7,300 km
7995	Ford	F-Series Pickup	2023	575000.0	35 km
7993	Toyota	Alphard	2024	550000.0	0 km
7994	Toyota	Alphard	2024	550000.0	0 km
7992	Toyota	Alphard	2024	540000.0	0 km
7991	Ford	F-Series Pickup	2023	539000.0	1,600 km
7990	Toyota	Alphard	2024	530000.0	0 km
7989	Toyota	Alphard	2024	520000.0	0 km
7988	Toyota	Alphard	2024	515000.0	2,200 km
7987	Ford	Shelby Cobra	2015	499000.0	5,000 km
7986	Ford	Shelby Cobra	1965	499000.0	23,433 km



In [25]:

```
# Calculate and display some statistics
print("\n
Price Statistics:")
print(df['Price'].describe())

print("\n
Top 5 Most Expensive Car Brands (Average Price):")
print(df.groupby('Car Brand')['Price'].mean().sort_values(ascending=False).head())

print("\n
Number of Cars by Brand:")
print(df['Car Brand'].value_counts().head(10))
```

Price Statistics:

count	8.006000e+03
mean	8.796613e+04
std	1.131822e+05
min	3.500000e+03
25%	3.242500e+04
50%	5.900000e+04
75%	1.120000e+05
max	4.750000e+06

Name: Price, dtype: float64

Top 5 Most Expensive Car Brands (Average Price):

Car Brand

Toyota	119731.127676
Ford	98059.913651
Chevrolet	79095.475610
Nissan	76283.303059
Kia	57978.809826

Name: Price, dtype: float64


Number of Cars by Brand:

Toyota	2522
Nissan	2125
Ford	1216
Hyundai	850
Kia	631
Honda	580
Chevrolet	82

Name: Car Brand, dtype: int64



```
In [27]: # Create a box plot of prices by car brand
plt.figure(figsize=(15, 8))
sns.boxplot(x='Car Brand', y='Price', data=df, fliersize=1)
plt.title('Price Distribution by Car Brand')
plt.xticks(rotation=90)
plt.ylabel('Price (AED)')
plt.yscale('log') # Using log scale for better visualization
plt.tight_layout()
plt.show()
```



The Ford GT models dominate the top spots with prices ranging from 2,649,000 to 4,750,000 AED. The average price of cars in the dataset is approximately 87,966 AED, with a standard deviation of 113,182 AED. Toyota, Ford, and Chevrolet are among the top brands with the highest average prices. Toyota also has the highest number of cars listed.



```
In [28]: # Calculate average price for each brand
brand_avg_price = df.groupby('Car Brand')['Price'].mean().sort_values(ascending=False)

# Get top 20 brands
top_20_brands = brand_avg_price.head(20)

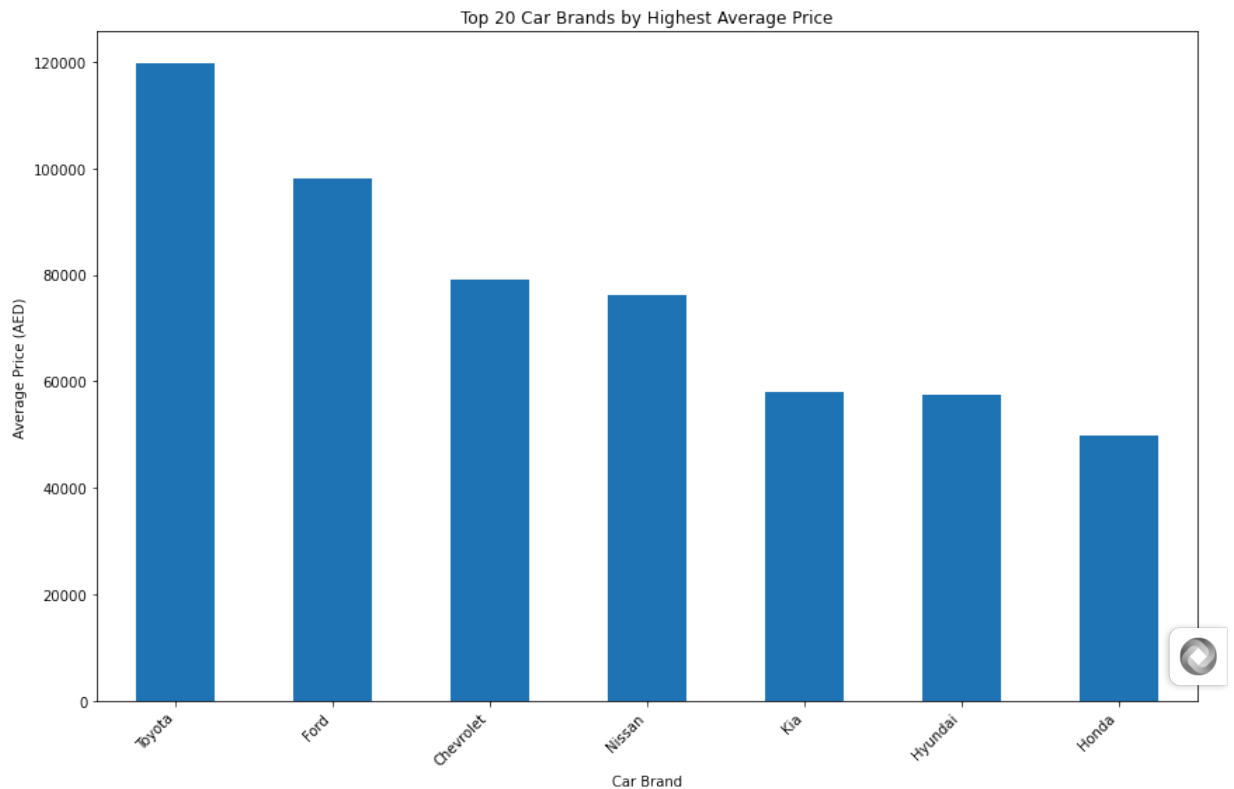
print("Top 20 Car Brands by Highest Average Price:")
print(top_20_brands)

# Create a bar plot
plt.figure(figsize=(12, 8))
top_20_brands.plot(kind='bar')
plt.title('Top 20 Car Brands by Highest Average Price')
plt.xlabel('Car Brand')
plt.ylabel('Average Price (AED)')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

Top 20 Car Brands by Highest Average Price:

Car Brand	Average Price (AED)
Toyota	119731.127676
Ford	98059.913651
Chevrolet	79095.475610
Nissan	76283.303059
Kia	57978.809826
Hyundai	57566.777647
Honda	49913.512069

Name: Price, dtype: float64



```
In [2]: # Step 2: Data Cleaning
# Convert Mileage and Price to numeric values
df['Mileage'] = df['Mileage'].str.replace(' km', '').str.replace(',', '').astype(float)
df['Price'] = df['Price'].str.replace(',', '').astype(float)

# Handle missing values (if any)
df = df.dropna()

# Encode categorical variables
df['Brand_Encoded'] = df['Car Brand'].astype('category').cat.codes
df['Model_Encoded'] = df['Car Model'].astype('category').cat.codes
df['Specs_Encoded'] = df['Specs'].astype('category').cat.codes

# Display the cleaned dataframe
print(df.head())
```

	Car Brand	Car Model	Production Year	Mileage	Price \
0	Nissan	Altima	2005	445740.0	3500.0
1	Toyota	Camry	1999	200000.0	5500.0
2	Ford	Focus	2006	366135.0	5500.0
3	Toyota	Echo	2005	200000.0	6000.0
4	Chevrolet	Epica	2009	250000.0	6000.0

	Description	Specs \
0	Dubai	GCC Specs
1	Perfect Condition Toyota Camry	GCC Specs
2	FORD FOCUS	GCC Specs
3	GCC - TOYOTA ECHO 2005 - Manual, Urgent Sale	GCC Specs
4	Chevrolet Epica	American Specs

	Timestamp	Location	Brand_Encoded	Model_Encoded	Specs_Encoded
0	04-03-24 14:49	Dubai	5	8	4
1	04-03-24 14:49	Dubai	6	24	4
2	04-03-24 14:49	Dubai	1	60	4
3	04-03-24 14:49	Dubai	6	44	4
4	45354.94097	Abu Dhabi	0	48	0




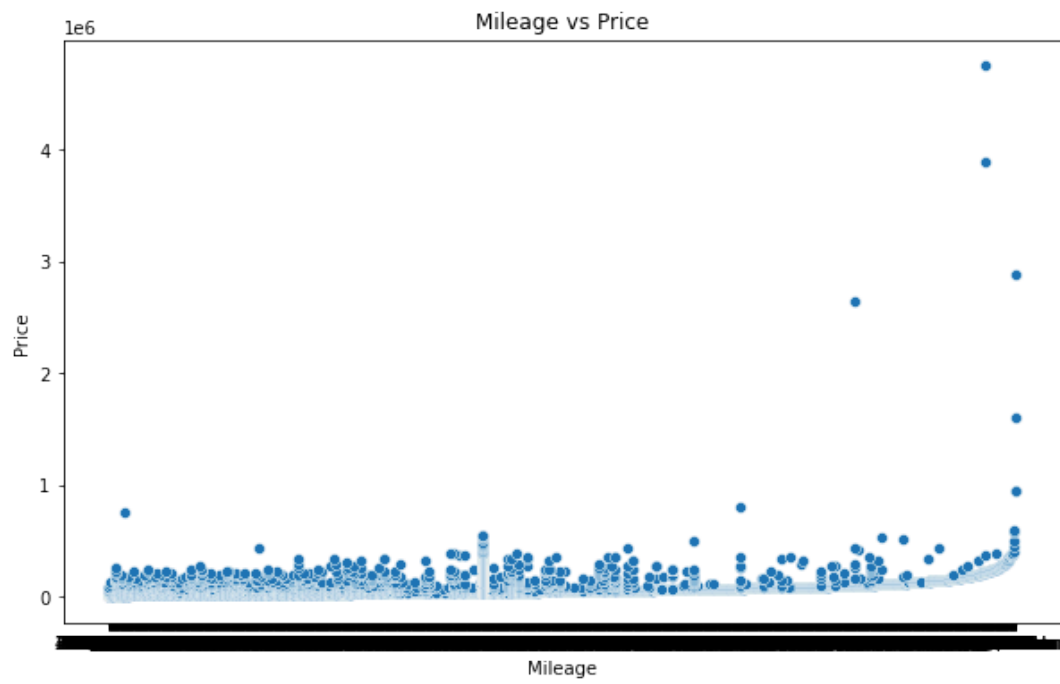

```
In [16]: # Correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()

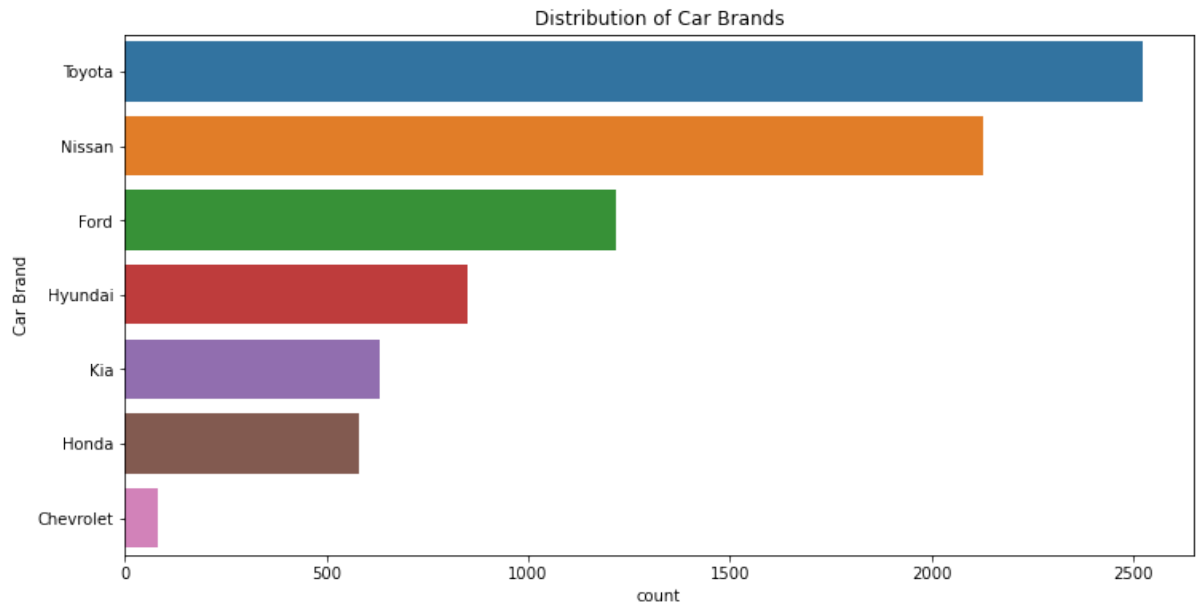
# Scatter plot: Production Year vs Price
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Production Year', y='Price', data=df)
plt.title('Production Year vs Price')
plt.show()

# Scatter plot: Mileage vs Price
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Mileage', y='Price', data=df)
plt.title('Mileage vs Price')
plt.show()


# Distribution of car brands
plt.figure(figsize=(12, 6))
sns.countplot(y='Car Brand', data=df, order=df['Car Brand'].value_counts().index)
plt.title('Distribution of Car Brands')
plt.show()
```







```
In [6]: # Distribution of car models
plt.figure(figsize=(12, 36))
sns.countplot(y='Car Model', data=df, order=df['Car Model'].value_counts().index)
plt.title('Distribution of Car Models')
plt.show()
```



```
In [7]: # Step 4: Feature Engineering
# Select features and target variable
features = ['Production Year', 'Mileage', 'Brand_Encoded', 'Model_Encoded', 'Specs_Encoded']
target = 'Price'

X = df[features]
y = df[target]

# Split the data into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [10]: # Step 5: Modeling
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Initialize and train the model
model = RandomForestRegressor(random_state=42)
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print('Mean Squared Error:', mse)
print('R-squared Score:', r2)

# Feature importance
feature_importance = pd.DataFrame({'feature': features, 'importance': model.feature_importances_})
print('Feature Importance:')
print(feature_importance.sort_values(by='importance', ascending=False))
```

Mean Squared Error: 2467218384.11639

R-squared Score: 0.632873800374377

Feature Importance:

	feature	importance
3	Model_Encoded	0.428705
1	Mileage	0.296152
0	Production Year	0.166136
2	Brand_Encoded	0.064026
4	Specs_Encoded	0.044981


Mean Squared Error: [2467218384.11639](#)

R-squared Score: 0.632873800374377

The R-squared score of about 0.63 indicates that our model explains approximately 63% of the variance in car prices. This suggests that while our model has some predictive power, there are likely other factors influencing car prices that aren't captured in our dataset.



```
In [9]: ▶ # Step 6: Visualizations
# Actual vs Predicted Prices
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=2)
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Actual vs Predicted Prices')
plt.show()
```



This scatter plot compares the actual prices of cars in our test set with the prices predicted by our model. The diagonal line would represent perfect predictions. The spread of points around this line indicates the model's accuracy.

Based on this analysis, we can conclude that Mileage and Production Year are the most influential factors in determining a used car's price in this UAE dataset. The car's brand also plays a role, with luxury brands generally commanding higher prices. However, there are likely other factors not captured in this dataset that also influence car prices, as our model only explains about 63% of the price variance.

In []: ▶

